

Chapter 14 Report: Classifying Images with CNNs

Quan Pham

June 18, 2025

Contents

1	Abstract	2
2	Key Concepts	2
2.1	Convolutional Neural Networks	2
2.2	Discrete Convolution	2
2.3	Pooling Layers	2
2.4	Convolution layer	2
2.5	Regularization	3
2.6	Loss Functions	3
2.7	Other Techniques	3
3	Algorithms and Models Studied	3
4	Code Examples and Practical Exercises	3
4.1	Implementing a deep CNN using PyTorch	3
4.2	Smile classification from face images using CNN	4
5	Learnings and Challenges	6
5.1	Key Takeaways	6
5.2	Challenges Encountered	6
5.3	Opening Questions	6
6	Conclusion and Furture work	6
6.1	Conclusion	6
6.2	Future work	6

1 Abstract

In this chapter, I studied convolutional neural networks (CNNs) and their application in image classification tasks. CNNs are designed to extract hierarchical features from images, which makes them highly effective for vision-related problems.

2 Key Concepts

2.1 Convolutional Neural Networks

- CNNs mimic the human visual cortex and automatically learn features from raw data.
- Feature hierarchies: low-level features \rightarrow high-level representations.
- Key ideas: **sparse connectivity** and **parameter sharing**.
- Pooling layers reduce spatial dimensions and help with generalization.

2.2 Discrete Convolution

$$y[i] = \sum_{k=-\infty}^{+\infty} x[i-k] \cdot w[k]$$
$$Y[i, j] = \sum_{k_1} \sum_{k_2} X[i-k_1, j-k_2] \cdot W[k_1, k_2]$$

Important hyperparameters: padding (full, same, valid), stride.

2.3 Pooling Layers

- Max-pooling and mean-pooling.
- Improve robustness to noise, reduce overfitting.
- Can use overlapping or non-overlapping pooling.

2.4 Convolution layer

- Handling multiple input/output channels:

$$Z^{\text{conv}}[:, :, k] = \sum_{c=1}^{C_{in}} W[:, :, c, k] * X[:, :, c]$$
$$Z[:, :, k] = Z^{\text{conv}} + b[k]$$
$$A[:, :, k] = \sigma(Z[:, :, k])$$

2.5 Regularization

- L2 regularization (weight decay) and dropout help prevent overfitting.
- Dropout encourages robust feature learning.

2.6 Loss Functions

- Binary classification: BCEWithLogitsLoss, BCELoss
- Multiclass: CrossEntropyLoss, NLLLoss

2.7 Other Techniques

- **Data augmentation:** enhances generalization.
- **Global average pooling:** reduces parameters.

3 Algorithms and Models Studied

4 Code Examples and Practical Exercises

4.1 Implementing a deep CNN using PyTorch

Prepare the MNIST dataset:

```
1 import torchvision
2 from torchvision import transforms
3 import torch
4 x = torch.tensor([5, 5, 5])
5
6 from torch import nn
7
8 image_path = '../data/'
9 transform = transforms.Compose([
10     transforms.ToTensor()
11 ])
12 mnist_dataset = torchvision.datasets.MNIST(root=image_path, train=True,
13     ↪ transform=transform, download=False)
14
15 from torch.utils.data import Subset
16 mnist_valid_dataset = Subset(mnist_dataset, torch.arange(10000))
17 mnist_train_dataset = Subset(mnist_dataset, torch.arange(10000,
18     ↪ len(mnist_dataset)))
19 mnist_test_dataset = torchvision.datasets.MNIST(root=image_path,
20     ↪ train=False, transform=transform, download=False)
```

```

18
19 from torch.utils.data import DataLoader
20 batch_size = 64
21 torch.manual_seed(1)
22 train_dl = DataLoader(mnist_train_dataset, batch_size, shuffle=True)
23 valid_dl = DataLoader(mnist_valid_dataset, batch_size, shuffle=False)

```

Construct the CNN model using the Sequential class:

```

1 model = nn.Sequential()
2 model.add_module(
3     'conv1',
4     nn.Conv2d(
5         in_channels=1, out_channels=32, kernel_size=5, padding=2
6     )
7 )
8 model.add_module('relu1', nn.ReLU())
9 model.add_module('pool1', nn.MaxPool2d(kernel_size=2))
10 model.add_module(
11     'conv2',
12     nn.Conv2d(
13         in_channels=32, out_channels=64, kernel_size=5, padding=2
14     )
15 )
16 model.add_module('relu2', nn.ReLU())
17 model.add_module('pool2', nn.MaxPool2d(kernel_size=2))
18 model.add_module('fc1', nn.Linear(3136, 1024))
19 model.add_module('relu3', nn.ReLU())
20 model.add_module('dropout', nn.Dropout(p=0.5))
21 model.add_module('fc2', nn.Linear(1024, 10))

```

Train accuracy and validation accuracy for 20 epochs:

```

Epoch 1 accuracy: 0.9496 val_accuracy: 0.9792
...
Epoch 10 accuracy: 0.9964 val_accuracy: 0.9902
...
Epoch 20 accuracy: 0.9981 val_accuracy: 0.9909

```

4.2 Smile classification from face images using CNN

```

1 get_smile = lambda attr: attr[18]
2
3 transform_train = transforms.Compose([
4     transforms.RandomCrop([178, 178]),

```

```

5     transforms.RandomHorizontalFlip(),
6     transforms.Resize([64, 64]),
7     transforms.ToTensor(),
8 ])
9
10 transform = transforms.Compose([
11     transforms.CenterCrop([178, 178]),
12     transforms.Resize([64, 64]),
13     transforms.ToTensor(),
14 ])

```

```

1 celeba_train_dataset = torchvision.datasets.CelebA(image_path,
  ↳ split='train', target_type='attr', download=False,
  ↳ transform=transform_train, target_transform=get_smile)
2 celeba_train_dataset = Subset(celeba_train_dataset, torch.arange(16000))
3 celeba_valid_dataset = Subset(celeba_valid_dataset, torch.arange(1000))

```

```

1 batch_size = 32
2 torch.manual_seed(1)
3 train_dl = DataLoader(celeba_train_dataset, batch_size, shuffle=True)
4 valid_dl = DataLoader(celeba_valid_dataset, batch_size, shuffle=False)
5 test_dl = DataLoader(celeba_test_dataset, batch_size, shuffle=False)

```

```

1 model = nn.Sequential()
2 model.add_module('conv1', nn.Conv2d(in_channels=3, out_channels=32,
  ↳ kernel_size=3, padding=1))
3 model.add_module('relu1', nn.ReLU())
4 model.add_module('pool1', nn.MaxPool2d(kernel_size=2))
5
6 model.add_module('conv2', nn.Conv2d(in_channels=32, out_channels=64,
  ↳ kernel_size=3, padding=1))
7 model.add_module('relu2', nn.ReLU())
8 model.add_module('pool2', nn.MaxPool2d(kernel_size=2))
9
10 model.add_module('conv3', nn.Conv2d(in_channels=64, out_channels=128,
  ↳ kernel_size=3, padding=1))
11 model.add_module('relu3', nn.ReLU())
12 model.add_module('pool3', nn.MaxPool2d(kernel_size=2))
13
14 model.add_module('conv4', nn.Conv2d(in_channels=128, out_channels=256,
  ↳ kernel_size=3, padding=1))
15 model.add_module('relu4', nn.ReLU())
16
17 model.add_module('pool4', nn.AvgPool2d(kernel_size=8))

```

```
18 model.add_module('flatten', nn.Flatten())
19 model.add_module('fc', nn.Linear(256, 1))
20 model.add_module('sigmoid', nn.Sigmoid())
```

```
1 loss_fn = nn.BCELoss()
2 optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
```

```
Epoch 1 accuracy: 0.6350 val_accuracy: 0.7100
...
Epoch 15 accuracy: 0.8899 val_accuracy: 0.8990
...
Epoch 30 accuracy: 0.9139 val_accuracy: 0.9030
```

5 Learnings and Challenges

5.1 Key Takeaways

5.2 Challenges Encountered

5.3 Opening Questions

6 Conclusion and Furture work

6.1 Conclusion

6.2 Future work

References

- [1] Yann LeCun, Bernhard Boser, John Denker, Donnie Henderson, Richard Howard, Wayne Hubbard, and Lawrence Jackel. Handwritten digit recognition with a back-propagation network. *Advances in neural information processing systems*, 2, 1989.